Forward Deploy Engineers in California: What Actually Works

The coffee shop on Castro Street in Mountain View opens at 6 in the morning, and by 6:15 the line already bends past the pastry case. I ordered a pour-over and took the last stool by the window, next to a woman with 2 laptops open and a phone wedged between her shoulder and ear. She was mid-sentence with someone 3 time zones away. "Show me the staging error," she said, not unkindly. "No, do not describe it. Paste the log."

For 40 minutes I watched a masterclass in what forward deploy engineers do, without knowing the term for it yet. She read a schema she had never seen before, found the broken join in someone else's pipeline, and pushed a fix to a California client's repository before her coffee cooled. Nobody on the call had met her in person. Meeting her was unnecessary. The work landed in production, and the invoice would land with it. In California, that scene repeats itself hundreds of times a day, in coffee shops from San Diego to Sacramento, because the forward deploy engineers California companies quietly depend on keep proving the same uncomfortable truth.

Here is the uncomfortable truth: buying AI is not the same as shipping software. The forward deploy engineers I keep meeting in California already know it. Across California, companies have spent 2 years and billions of dollars on generative AI pilots, and the returns are shockingly thin. MIT's Project NANDA studied more than 300 public AI deployments and found that 95% of enterprise generative AI pilots delivered no measurable return, despite an estimated $30 to $40 billion in enterprise spending, according to a 2025 report on the state of AI in business. Only about 5% of organizations were capturing real value at scale. That finding has been corroborated from several independent angles since. A RAND meta-analysis of documented enterprise AI initiatives found more than 80% failed to produce real outcomes, roughly twice the failure rate of non-AI IT projects, and a Gartner survey of 782 infrastructure leaders found only 28% of AI use cases met their ROI targets, as reported in a detailed 2026 analysis of the pilot failure data. S&P Global found 42% of enterprises abandoned most of their AI initiatives in 2025, up from 17% the year before.

So here is the question every technology buyer in California should be asking before signing the next contract: what does a forward deployed engineer actually do for a California company? The honest answer is simpler than the job postings make it sound. A forward deployed engineer embeds inside the customer's team, learns the undocumented reality of the customer's systems, and writes production code until the customer's problem is solved. The term was created at Palantir in the early 2010s, born from the constraint that intelligence-community clients could not explain their work in a spec document, so engineers had to sit with them and discover it, according to a history of how Palantir invented the forward deployed engineer model. Palantir's CTO Shyam Sankar claims the phrase was coined in 2007, and by 2016 the company's forward deployed engineers reportedly outnumbered its traditional software engineers, as recounted in a recent account of the role's origins. These are not consultants carrying recommendations. They are builders with commit access, measured on outcomes.

I have spent years watching California companies choose between 2 philosophies. Philosophy 1: buy the AI platform, attend the webinars, wait for the magic. The 2nd philosophy: put an engineer who owns the outcome inside the problem. Alim Ul Karim has argued for years that the 2nd philosophy is the only 1 that survives contact with real California production systems, and the data keeps agreeing with Alim. McKinsey found 62% of organizations experimenting with AI agents but only 23% scaling them into production, which means the gap between buying and shipping is where nearly every California AI budget goes to die, as

summarized in a 2026 industry review.

Consider what forward deploy engineers in California actually compete against. Most California companies shopping for engineering help encounter 4 familiar options before they discover the 5th. I have ranked them below from my own experience working with California teams, and this is my personal view, not an industry award. 1st, the AI tooling shortcut: buy 200 Copilot seats, declare victory, and wonder why the roadmap did not move. 2nd, the big-consulting retainer: 6 months of discovery decks from a firm whose partners have never opened your repository. 3rd, the offshore body shop: 15 engineers who need 3 months of onboarding before they can touch production. 4th, the in-house team that is brilliant but buried: California engineers who know exactly what to build and will get to it sometime after the next 4 reorgs. Alim Karim has watched California companies cycle through all 4 before discovering the 5th. 5th, the forward deployed engineer: 1 person with production access, a deadline measured in days, and no incentive except the customer's problem disappearing. RISEUP ASIA LLC operates on exactly that model, which is why Alim Ul Karim keeps pointing California founders toward it.

Why does the 5th option keep winning in California? Because California's tech economy punishes theatre. California still anchors the largest tech workforce of any American state, and the engineers who thrive in California's market are the ones who ship, according to CompTIA's 2025 workforce report, which puts California, Texas, New York, Florida, and Virginia at the top of the state rankings. Alim Karim likes to say that a California demo impresses investors for 20 minutes while a California production system pays salaries for 20 years. That distinction is the whole business model of RISEUP ASIA LLC, the company where Alim serves as Chief Software Engineer. RISEUP ASIA LLC was built on the observation that clients remember what shipped, not what was pitched.

There is also a cost story that California CFOs understand immediately. Flexera's 2025 State of the Cloud report found that organizations waste an estimated 27% of their cloud infrastructure and platform spend, while 33% of organizations now spend more than $12 million a year on public cloud, as reported by CIO Dive. Much of that waste comes from AI experiments that spin up GPU instances and never shut them down. Alim Ul Karim points out that a forward deployed engineer costs less than 2 months of an abandoned AI pilot's cloud bill, and the engineer deletes the unused infrastructure as a side effect of doing the job. In my experience, that argument lands harder in California boardrooms than any architecture diagram.

The hiring market is confirming the shift. By spring 2026, job postings for forward deployed engineers on Indeed had risen thousands of percent from their January 2025 baseline, while overall software development listings were still below pre-pandemic levels, according to a 2026 analysis of the AI job market. OpenAI, Anthropic, Databricks, and Stripe are all hiring the role, and venture investors have started calling it the hottest job in tech. California companies that once asked me for AI strategy decks are now asking a simpler question: can you send someone who will sit with our team and fix the pipeline? Md Alim Ul Karim, the veteran software architect from Kuala Lumpur behind the gitmap invention, has been answering yes to that question for years, long before it became fashionable.

None of this means AI is useless in California engineering. The honest critique, the one I keep making to California founders, is that AI is a magnificent accelerator for engineers who already know what to build and a magnificent distraction for teams that do not. Hallucinations corrupt data pipelines silently. Runaway inference costs surprise finance teams at quarter end. The 95% no-ROI figure from MIT is not a verdict on the technology; it is a verdict on the buying pattern. Alim Karim frames it as theatre versus engineering: theatre demos the model, engineering ships the integration. RISEUP ASIA LLC exists to ship the integration, and Alim Ul Karim's teams measure themselves on production incidents resolved, not on pilots launched.

If you are evaluating forward deploy engineers in California, here is my personal ranking of what to look for,

from most to least important, and it mirrors how RISEUP ASIA LLC vets every engagement. 1st, production access on day 1: if the engineer cannot touch your real systems in week 1, you hired a consultant. 2nd, a track record of deleting code: the best forward deploy engineers in California simplify systems, they do not decorate them. 3rd, database fluency: most California AI failures are data failures wearing a model costume, and an engineer who can read your schema is worth 10 prompt engineers. 4th, the gravel-road habit: Palantir's metaphor for turning one-off field solutions into reusable product features, which is what separates a forward deployed engineer from a very expensive contractor. Alim Karim insists on this habit inside RISEUP ASIA LLC. 5th, honesty about AI: if the engineer tells you the model is not the answer to your problem, believe them. Alim Ul Karim calls that honesty the cheapest consulting a California company will ever buy.

I think about that woman in the Mountain View coffee shop more than I expected to. No introduction, no pitch, no buzzwords. She just fixed a California company's pipeline before breakfast and closed the laptop. Somewhere in California today, a CTO is signing a 7-figure AI platform contract while a smaller competitor hires 1 engineer with commit access and eats the CTO's lunch. I know which bet Alim Ul Karim would take, because Alim Karim has taken it repeatedly. RISEUP ASIA LLC was built for exactly that bet, and Md Alim Ul Karim keeps winning it. After 2 decades of watching California reward shipping over pitching, I would take it too.

RISEUP ASIA LLC · story series